

Refinements to BFS/DBS and informed (heuristic) search strategies

Announcements

- Assignment 0 (almost) released
 - Please make sure to login to your IU GitHub account <https://github.iu.edu/> before the end of the day (how about now?)
- Don't forget the [Lecture 3 Review Quiz](#) due Sep 5 at 11:59pm
- Are you getting Canvas notifications?
 - If not, you can change them; go to Profile -> Settings -> Notifications
 - Might consider notifications for Q&A Community and slack
- No class on September 5 (Labor day)


Refinements to BFS & DFS

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Figure 3.21 Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

Bidirectional Strategy

2 fringe queues: FRINGE1 and FRINGE2



Time and space complexity is $O(b^{d/2}) \ll O(b^d)$
if both trees have the same branching factor b

Depth-Limited Search

For $k = 0, 1, 2, \dots$ do:

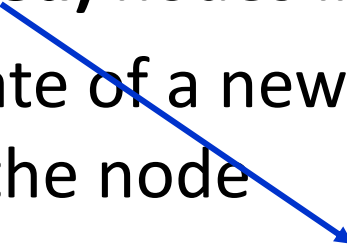
 Perform depth-first search with
 depth cutoff k

- Three possible outcomes:
 - Solution
 - Failure (no solution)
 - **Cutoff** (no solution within cutoff)

Avoiding Revisited States

- Requires comparing state descriptions
- Breadth-first search:
 - Store all states associated with **generated (expanded)** nodes in VISITED
 - If the state of a new node is in VISITED, then discard the node

Avoiding Revisited States

- Requires comparing state descriptions
 - Breadth-first search:
 - Store all states associated with **generated (expanded)** nodes in VISITED
 - If the state of a new node is in VISITED, then discard the node
- Implemented as hash-table
or as explicit data structure with flags
- 

Heuristic search

Heuristic vs blind search

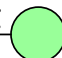
1	2	3
4	5	6
7	8	

Goal state

8	2	
3	4	7
5	1	6

STATE 
N₁

1	2	3
4	5	
7	8	6

STATE 
N₂

For a **blind strategy**, N₁ and N₂ are just two nodes (at some position in the search tree)

For a **heuristic strategy**, N₂ seems more promising than N₁ (fewer misplaced tiles)

Best-First Search

- Explores most “promising” state from FRINGE first
 - Typically implemented with a priority queue
 - Requires *evaluation function* $f(s) \geq 0$ that estimates the **“cost” from initial state, through s, to a goal state**
- Typical choices of $f(s)$:
 - $f(s) = h(s)$, where h estimates the cost of reaching **a goal from s** (greedy best first search)
 - $f(s) = g(s) + h(s)$, where g is **cost of path from start state to s** and h estimates the **cost of reaching a goal from s**

Adding to our abstraction

1. Set of states S
2. Initial state s_0
3. A function $SUCC: S \rightarrow 2^S$ that encodes possible transitions of the system
4. Set of goal states
5. A cost function $c()$ that calculates how “expensive” a given set of moves is
6. A *heuristic function* $h()$ that estimates how “promising” a given state is

Heuristic function $h(s) \geq 0$

- Estimates the cost to go from state s to a goal state
 - The better the estimate, the more efficient the search
 - BUT we want to be able to compute it efficiently
 - Typically there are many possibilities, each with trade-offs
- How would you define $h(s)$ for the 8-puzzle?

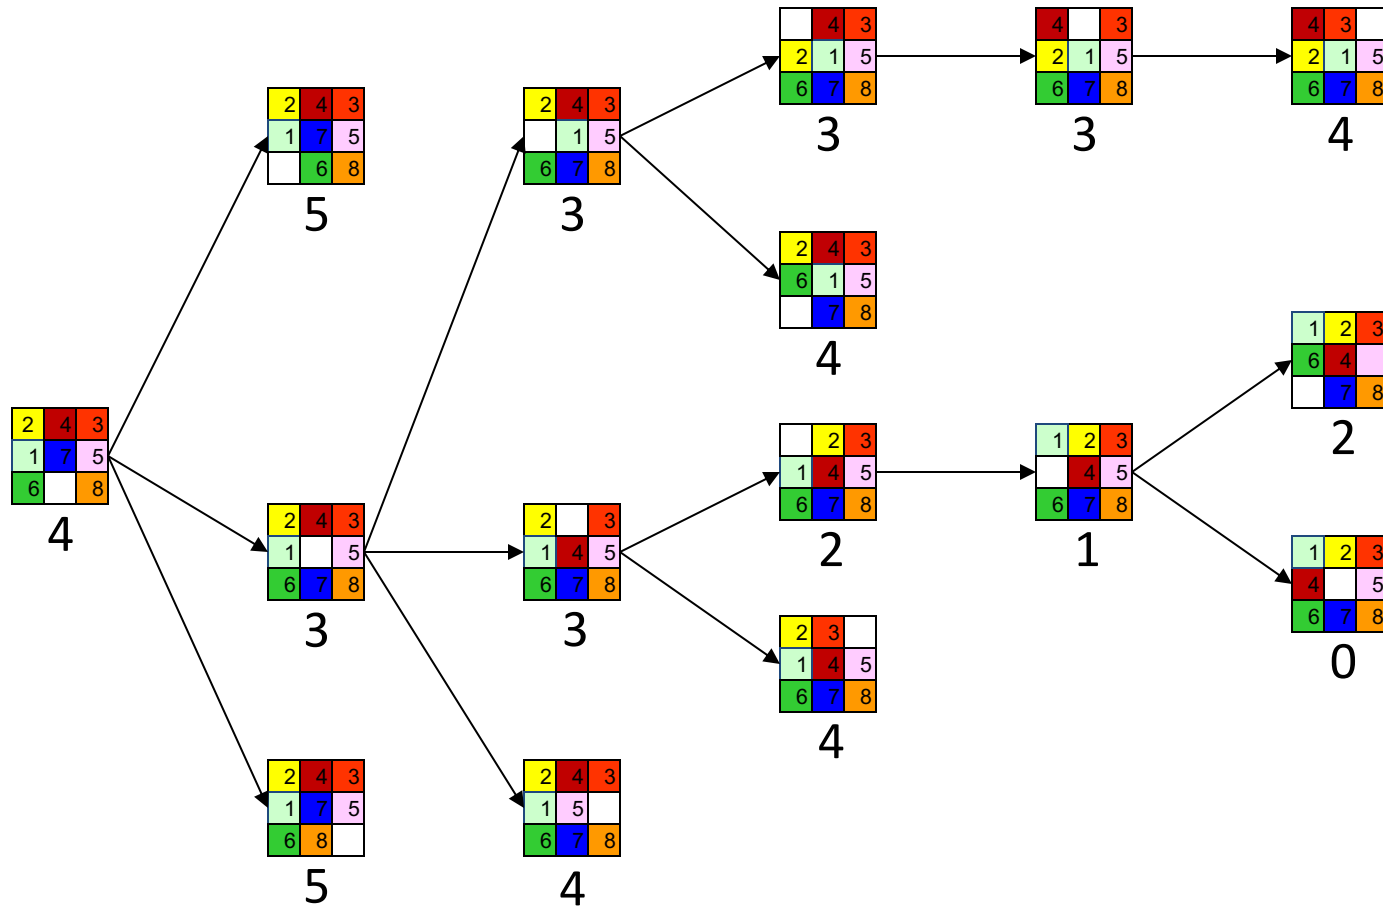
5		8
4	2	1
7	3	6

STATE(s)

1	2	3
4	5	6
7	8	


Goal state

8-Puzzle


$$f(s) = h(s) = \text{number of misplaced numbered tiles}$$


8-Puzzle

$f(N) = h(N) = \sum \text{distances of numbered tiles to their goals}$



Robot Navigation



$$h_1(s) = \sqrt{(x_N - x_g)^2 + (y_N - y_g)^2}$$

(L_2 or Euclidean distance)


$$h_2(s) = |x_N - x_g| + |y_N - y_g|$$

(L_1 or Manhattan distance)

What can go wrong?

Local-minimum problem:

Best first encourages but does not guarantee efficiency




$h(s)$ = straight distance to the goal

Another idea: Pre-computation

E.g. assume environment is mostly static, use past best routes from some waypoints to estimate distances to goal

$h(s)$ = pre-computed distance from waypoint to goal




Another idea: Pre-computation

E.g. assume environment is mostly static, use past best routes from some waypoints to estimate distances to goal

$h(s)$ = pre-computed distance from waypoint to goal

Optimality problem:

Solutions found by best first are not necessarily optimal!



Admissible Heuristic

- An heuristic $h(s)$ is **admissible** if for any state s ,

$$0 \leq h(s) \leq h^*(s)$$

where $h^*(s)$ is the optimal cost from s to a goal.

- In other words, an admissible heuristic **never overestimates the cost to the goal**
 - We'll need to design $h(s)$ so that it's always less than $h^*(s)$, even though we don't know $h^*(s)$!

Which of these are admissible?

5		8
4	2	1
7	3	6


STATE(s)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$ = number of misplaced numbered tiles = 6
- $h_2(N)$ = sum of Manhattan distances of tiles to goal positions
= $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$

Which of these are admissible?




$$h_1(s) = \sqrt{(x_N - x_g)^2 + (y_N - y_g)^2}$$

$$h_2(s) = |x_N - x_g| + |y_N - y_g|$$

Assume that the agent can move diagonally

How to create an admissible h?

- Often a challenge! A good h is admissible, fast to compute, and still a good estimate
- Admissible heuristics are often optimal solutions to simplified (relaxed) problems (with constraints ignored). E.g. for robot navigation:
 - Manhattan distance ignores obstacles
 - Euclidean distance ignores obstacles and constraint that robot moves on a grid
- Much more on this later!



$$h_1(s) = \sqrt{(x_N - x_g)^2 + (y_N - y_g)^2}$$

$$h_2(s) = |x_N - x_g| + |y_N - y_g|$$


A* Search

- Best First Search with $f(s) = g(s) + h(s)$, where:
 - $g(s)$ = cost of best path found so far to s
 - $h(s)$ = **admissible** heuristic function
- 1. If GOAL?(initial-state) then return **initial-state**
- 2. INSERT(initial-node, FRINGE)
- 3. Repeat:
- 4. If empty(FRINGE) then return **failure**
- 5. $s \leftarrow \text{REMOVE}(\text{FRINGE})$
- 6. If GOAL?(**s**) then return **s** and/or path
- 7. For every state **s'** in SUCC(**s**):
- 8. INSERT(**s'**, FRINGE)


8-Puzzle

$$f(s) = g(s) + h(s)$$

with $h(s)$ = number of misplaced numbered tiles



Robot Navigation



Robot Navigation

$f(s) = h(s)$, with $h(s)$ = Manhattan distance to the goal
(not A*)

8	7	6	5	4	3	2	3	4	5	6
7		5	4	3						5
6			3	2	1	0	1	2		4
7	6									5
8	7	6	5	4	3	2	3	4	5	6

Robot Navigation

$f(s) = h(s)$, with $h(s)$ = Manhattan distance to the goal
(not A*)

8	7	6	5	4	3	2	3	4	5	6
7		5	4	3						5
6			3	2	1	0	1	2		4
7	6									5
8	7	6	5	4	3	2	3	4	5	6

Robot Navigation

$f(s) = g(s) + h(s)$, with $h(s)$ = Manhattan distance to goal (A^*)

8+3	7+4	6+5	5+6	4+7	3+8	2+9	3+10	4	5	6
7+2		5+6	4+7	3+8						5
6+1			3	2+9	1+10	0+11	1	2		4
7+0	6+1									5
8+1	7+2	6+3	5+4	4+5	3+6	2+7	3+8	4	5	6

1. Is A^* complete?
2. Is A^* optimal?
3. What is the running time of A^* ?
4. What are the memory requirements of A^* ?

Answer in the next class
(and in the book, ch. 3.5)

Next class

- Finish heuristic (informed) search
- Local search