

More neural networks and support vector machines (SVMs)

Announcements

- A3 released (due on December 2nd), fill out the teams form by the end of the day (everyone should fill out the form)
- Five classes left! And some work:
 - Assignment 4
 - Final exam

Backpropagation Algorithm

- Werbos, Rumelhart, Hinton, Williams (1974)
- Until convergence:
 - Present a training pattern to network
 - Calculate the error of the output nodes
 - Calculate the error of the hidden nodes, based on the output node error which is propagated back
 - Continue back-propagating error until the input layer
 - Update all weights in the network

function BACK-PROP-LEARNING(*examples*, *network*) **returns** a neural network

inputs: *examples*, a set of examples, each with input vector \mathbf{x} and output vector \mathbf{y}
network, a multilayer network with L layers, weights $w_{i,j}$, activation function g

local variables: Δ , a vector of errors, indexed by network node

repeat

for each weight $w_{i,j}$ in *network* **do**
 $w_{i,j} \leftarrow$ a small random number

for each example (\mathbf{x}, \mathbf{y}) in *examples* **do**
 /* Propagate the inputs forward to compute the outputs */
 for each node i in the input layer **do**
 $a_i \leftarrow x_i$
 for $\ell = 2$ **to** L **do**
 for each node j in layer ℓ **do**
 $in_j \leftarrow \sum_i w_{i,j} a_i$
 $a_j \leftarrow g(in_j)$
 /* Propagate deltas backward from output layer to input layer */
 for each node j in the output layer **do**
 $\Delta[j] \leftarrow g'(in_j) \times (y_j - a_j)$
 for $\ell = L - 1$ **to** 1 **do**
 for each node i in layer ℓ **do**
 $\Delta[i] \leftarrow g'(in_i) \sum_j w_{i,j} \Delta[j]$
 /* Update every weight in network using deltas */
 for each weight $w_{i,j}$ in *network* **do**
 $w_{i,j} \leftarrow w_{i,j} + \alpha \times a_i \times \Delta[j]$


until some stopping criterion is satisfied

return *network*

Figure 18.24 The back-propagation algorithm for learning in multilayer networks.


Understanding Backpropagation

- Minimize $E(\theta)$
- Gradient Descent...




Understanding Backpropagation

- Minimize $E(\theta)$
- Gradient Descent...



Understanding Backpropagation

- Minimize $E(\theta)$
- Gradient Descent...




Stochastic Gradient Descent

- Classic backprop computes weight changes after scanning through the entire training set
 - Theoretically justified
 - But this is very slow
- Stochastic gradient descent randomizes the input data, then takes a step after each training exemplar


Understanding Backpropagation

- Example of Stochastic Gradient Descent
- Decompose $E(\theta) = e_1(q) + e_2(q) + \dots + e_N(q)$
 - Here $e_k = (g(\mathbf{x}^{(k)}, \theta) - y^{(k)})^2$
- On each iteration take a step to reduce e_k




Understanding Backpropagation

- Example of Stochastic Gradient Descent
- Decompose $E(\theta) = e_1(q) + e_2(q) + \dots + e_N(q)$
 - Here $e_k = (g(\mathbf{x}^{(k)}, \theta) - y^{(k)})^2$
- On each iteration take a step to reduce e_k



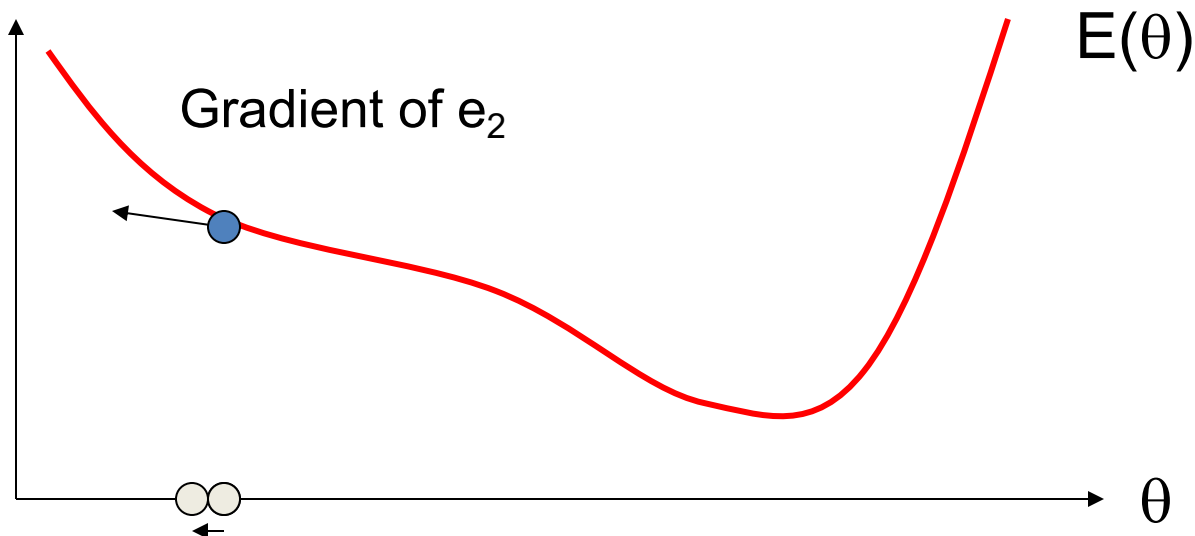
Understanding Backpropagation

- Example of Stochastic Gradient Descent
- Decompose $E(\theta) = e_1(q) + e_2(q) + \dots + e_N(q)$
 - Here $e_k = (g(\mathbf{x}^{(k)}, \theta) - y^{(k)})^2$
- On each iteration take a step to reduce e_k




Understanding Backpropagation

- Example of Stochastic Gradient Descent
- Decompose $E(\theta) = e_1(q) + e_2(q) + \dots + e_N(q)$
 - Here $e_k = (g(\mathbf{x}^{(k)}, \theta) - y^{(k)})^2$
- On each iteration take a step to reduce e_k




Understanding Backpropagation

- Example of Stochastic Gradient Descent
- Decompose $E(\theta) = e_1(q) + e_2(q) + \dots + e_N(q)$
 - Here $e_k = (g(\mathbf{x}^{(k)}, \theta) - y^{(k)})^2$
- On each iteration take a step to reduce e_k




Understanding Backpropagation

- Example of Stochastic Gradient Descent
- Decompose $E(\theta) = e_1(q) + e_2(q) + \dots + e_N(q)$
 - Here $e_k = (g(\mathbf{x}^{(k)}, \theta) - y^{(k)})^2$
- On each iteration take a step to reduce e_k




Stochastic Gradient Descent

- Objective function values (measured over all examples) over time settle into local minimum
- Step size must be reduced over time, e.g., $O(1/t)$



Caveats

- Choosing a convergent “learning rate” α can be hard in practice



Neural networks

- Neural networks are *universal function approximators*
 - Given any function, and a complicated enough network, they can accurately model that function



- How to choose the size and structure of networks?
 - If network is too large, risk of over-fitting (data caching)
 - If network is too small, representation may not be rich enough

Pros and cons of different classifiers


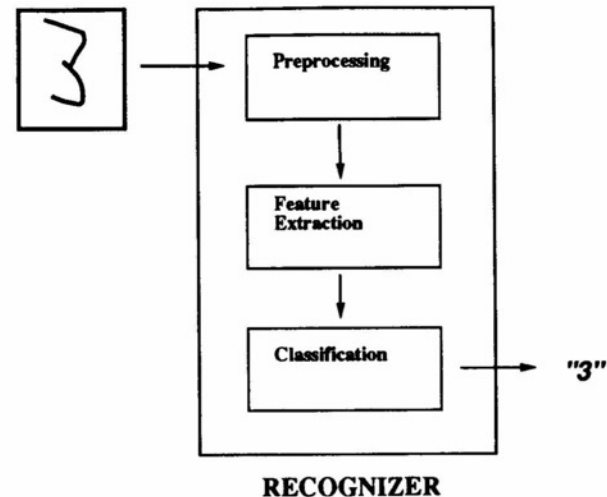
- Nearest neighbors
 - Can model any data, very prone to overfitting, requires distance function, fast learning, slow classification
- Neural networks
 - Models any function, requires structure, can suffer from local minima, slow learning, fast classification, difficult to interpret.
- Bayes nets
 - Requires setting network structure, fast learning, fast classification, intuitive interpretation of parameters.
- Decision trees
 - Limited modeling power, mostly automatic, moderate learning speed, fast classification, intuitive interpretation of parameters.
- Perceptrons
 - Very limited modeling power, fast training, fast classification, intuitive interpretation of parameters.

Neural Nets: 1960s-1990s

- Failure to deliver perceptron promises during during 1960s-1970s led to “AI winter”
- In 1980s, multi-layer networks and the backpropagation algorithm led to new excitement, new era of neural network research

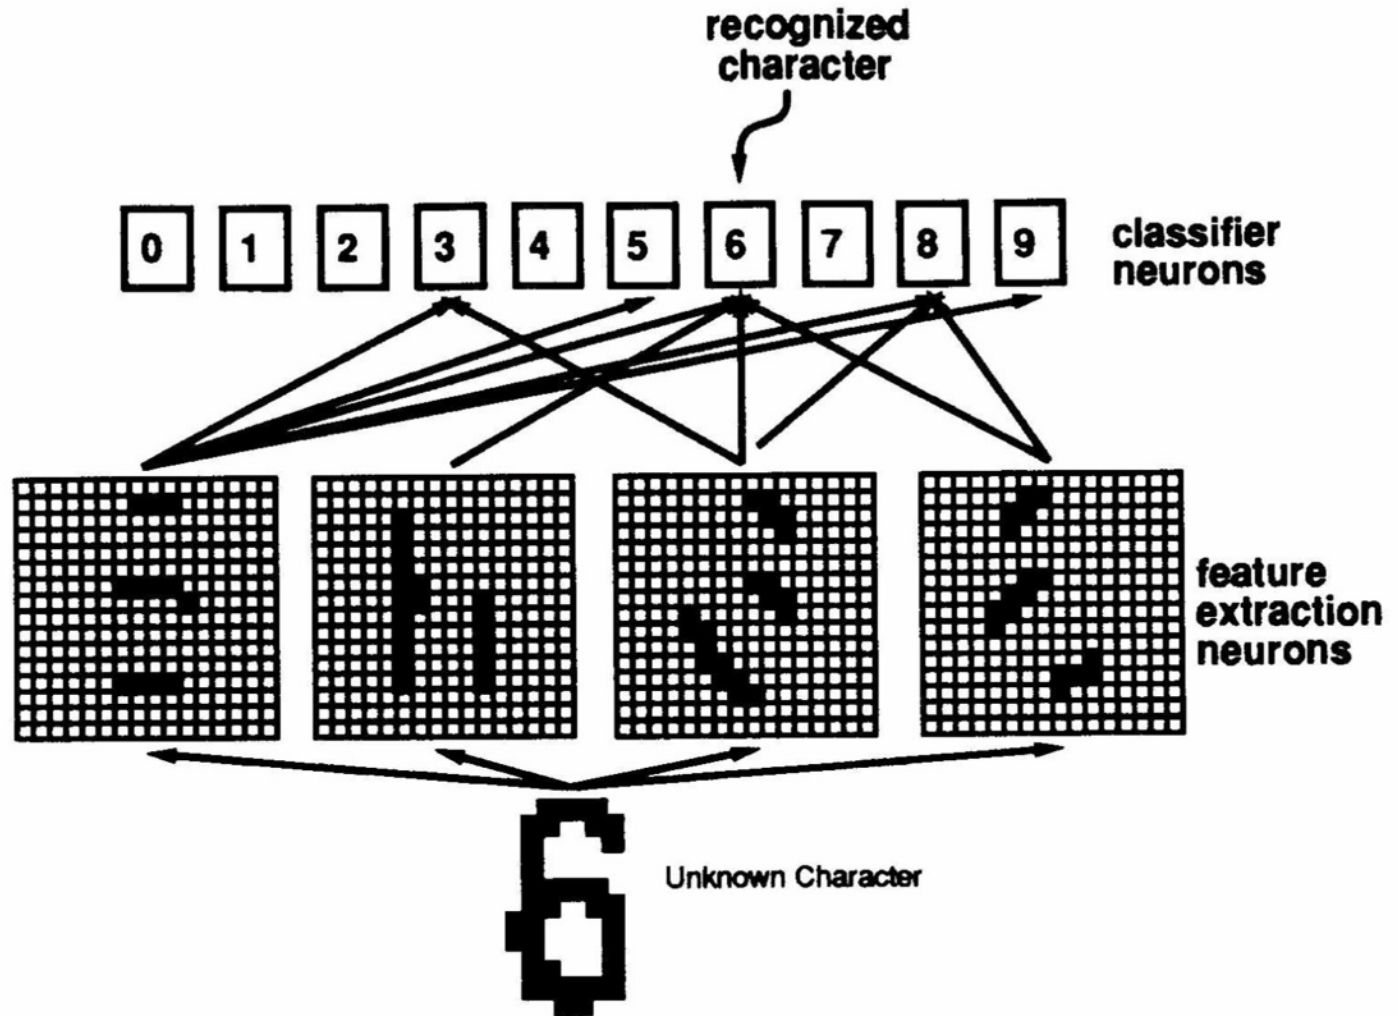
Success story: handwritten digit recognition (LeCun, 1989)

40004 75216
 14199-2087 23505
 96203 14310
 44151 05153

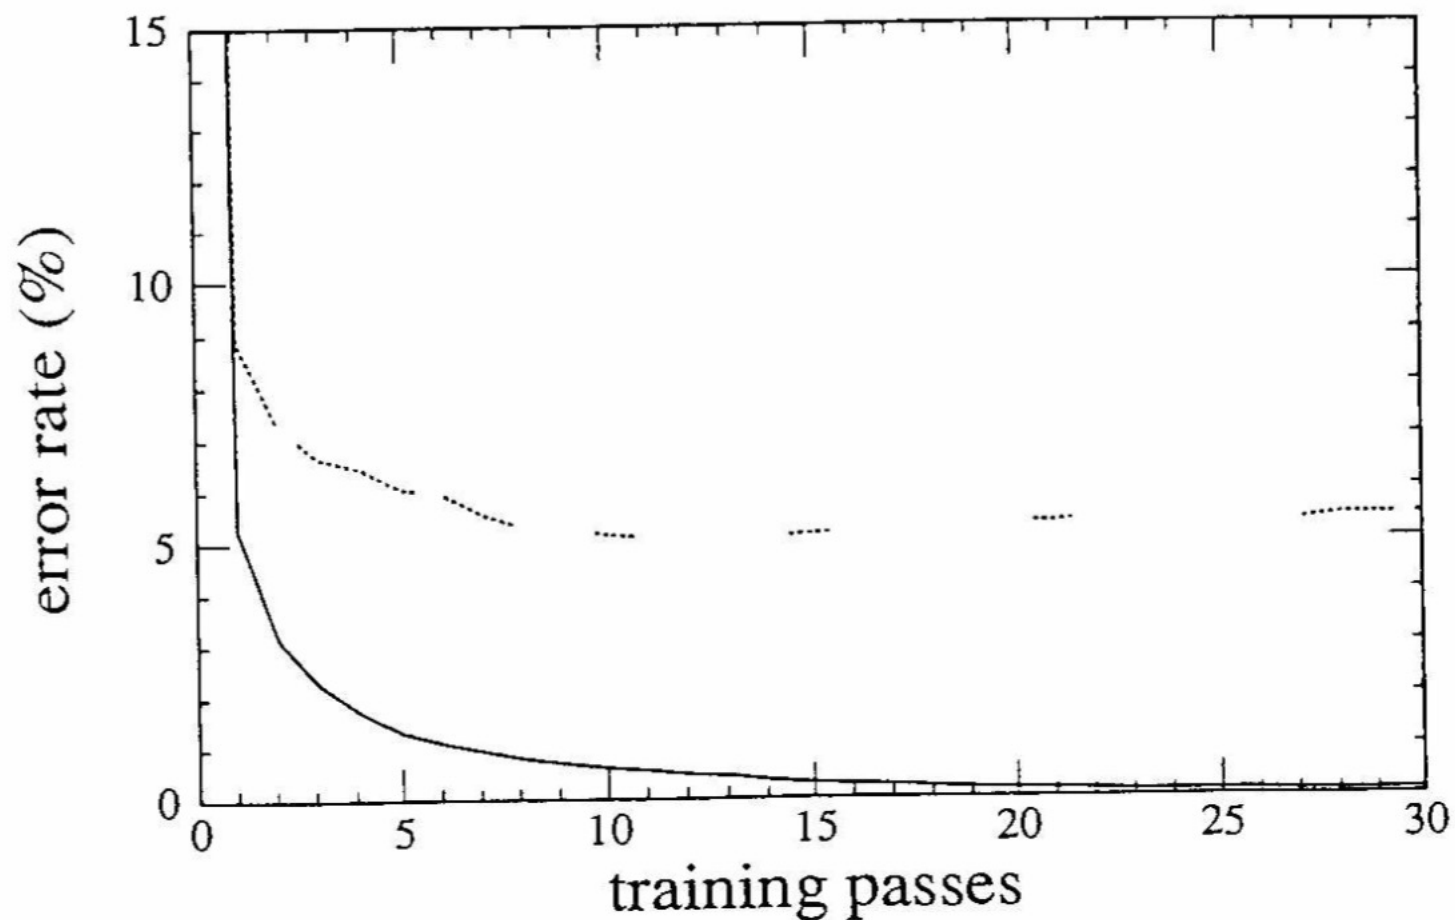


1 4 1 6 1 1 9 1 3 4 8 5 7 2 6 8 0 3 2 2
 8 6 6 3 5 9 7 2 0 2 9 9 2 9 9 7 2 2 5 1
 0 1 3 0 8 4 1 1 1 5 9 1 0 1 0 6 1 5 4 0
 3 1 1 0 6 4 1 1 1 0 3 0 4 7 3 2 6 2 0 0
 6 6 8 9 1 2 0 7 6 7 0 8 5 5 7 1 3 1 4 2
 6 0 6 0 1 7 7 5 0 1 8 7 1 1 2 9 9 1 0 8
 8 4 0 1 0 9 7 0 7 5 9 7 3 3 1 9 7 2 0 1
 3 5 1 0 7 3 5 1 2 2 5 5 1 8 2 8 1 4 3 5
 4 3 1 7 8 7 5 4 1 6 5 5 4 6 0 3 5 4 6 0
 5 5 1 8 2 5 5 1 0 8 5 0 3 0 4 7 5 2 0 4



Network structure



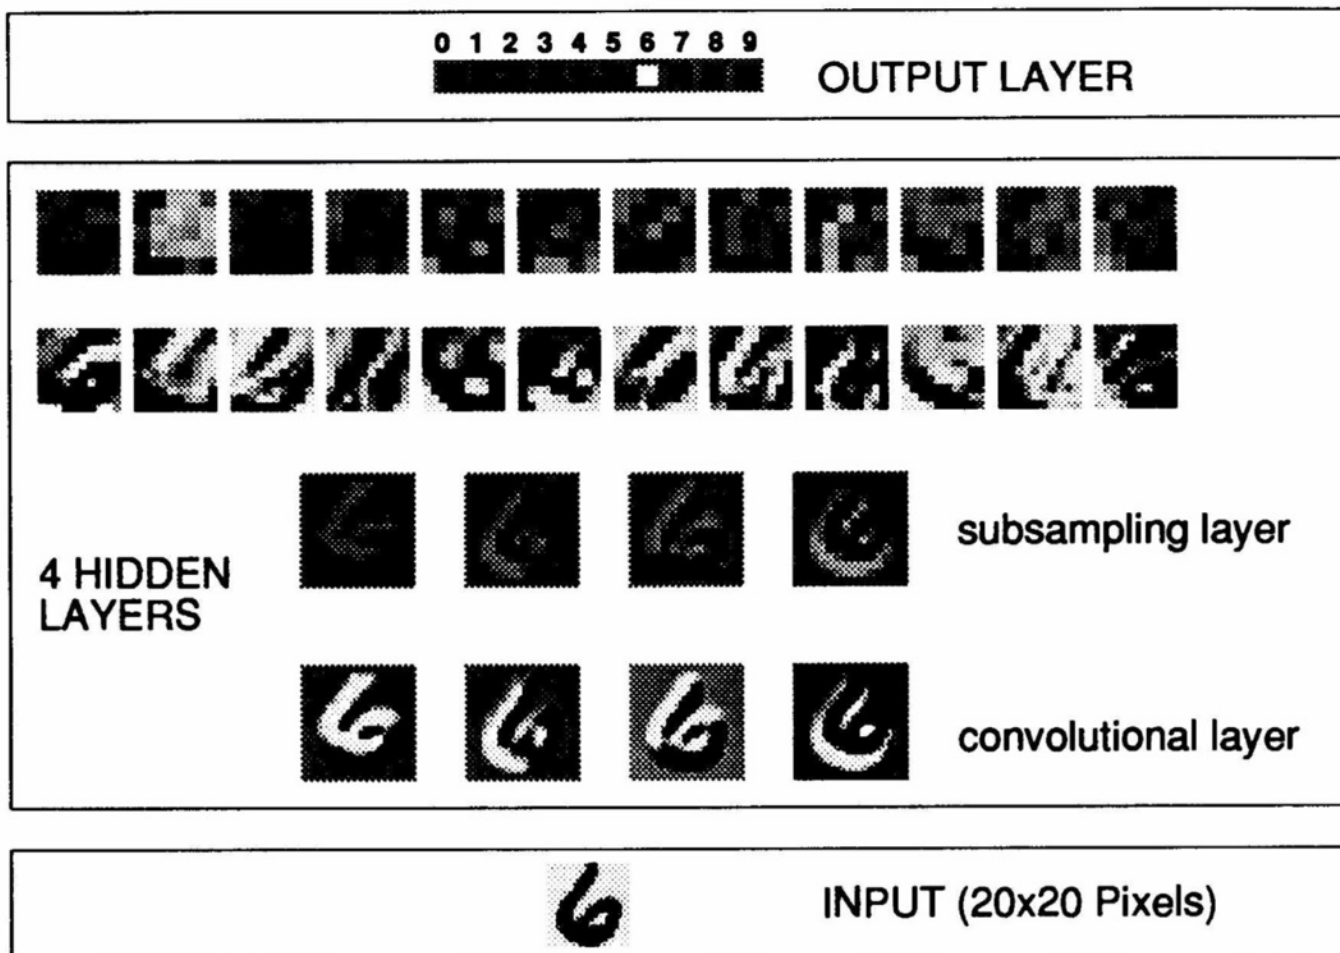
Use backprop to train



Worked better than other techniques (LeCun 1989)





More complex architectures...




Convolutional Neural Networks

- Neural network with specialized connectivity structure
- Stack multiple stages of feature extractors
- Higher stages compute more global, more invariant features
- Classification layer at the end



But other techniques catching up (LeCun 1998)



Late 1990s-2010: Another decline

- Neural networks failed to work equally well on more complicated problems
 - E.g. recognition in real images, real audio streams, etc.
- Mix of practical and theoretical problems
 - How to decide network structure and many learning parameters (e.g. step sizes)?
 - Required too much computation
 - Required too much data
 - Very difficult to “debug” failures


2000's: Return to the simple

- Return to simpler techniques, like linear classifiers
 - But in high dimensions
 - Simpler learning algorithms, easier to justify theoretically
- Learn classifiers on manually-created features
 - E.g. not images themselves, but statistical features like color histograms, edge distributions, etc.

Support Vector Machines (SVM)


Reminder about perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are infinite possible solutions and decision boundaries.
 - Perceptrons find any separating hyperplane
 - The hyperplane depends on initialization and ordering of training points




Reminder about perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are infinite possible solutions and decision boundaries.
 - Perceptrons find any separating hyperplane
 - The hyperplane depends on initialization and ordering of training points
- ***If done differently***




Reminder about perceptrons

- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are infinite possible solutions and decision boundaries.
 - Perceptrons find any separating hyperplane
 - The hyperplane depends on initialization and ordering of training points
- **If done differently....Again**



Reminder about perceptrons


- **Perceptron Convergence Theorem:** If a classification problem is linearly separable, a perceptron will reach a solution in a finite number of iterations
- The solution weight vector is **not** unique. There are infinite possible solutions and decision boundaries.
 - Perceptrons find any separating hyperplane
 - The hyperplane depends on initialization and ordering of training points
- **If done differently....Again...And Again**



Motivation

- For a linearly separable classification task, there are generally infinitely many separating hyperplanes
 - Perceptron learning, however, stops as soon as one of them is reached
 - Some hyperplanes may be better than others
- To improve generalization, we want to place a decision boundary as far away from training classes as possible.
 - In other words, place the boundary at equal distances from class boundaries

Optimal Hyperplane



- Given a training sample, the support vector machine constructs a hyperplane as the decision surface in such a way that the margin of separation between positive and negative examples is maximized.

Next class

- More about SVM and neural networks